



# Practical ML Advice

Rishabh Iyer

Based on slides from Jude Shavlik and Tom Dietterich

## Proper Experimental Methodology Can Have a Huge Impact:

A 2002 paper in *Nature* (a major journal) needed to be corrected due to “training on the testing set”


Original report : 95% accuracy (5% error rate)

Corrected report (which still is buggy):

73% accuracy (27% error rate)

Error rate increased over 400%!!!

# Some Typical ML Experiments



# Typical Experiments



	Test Set Performance
Full System	80%
Without Module A	75%
Without Module B	62%

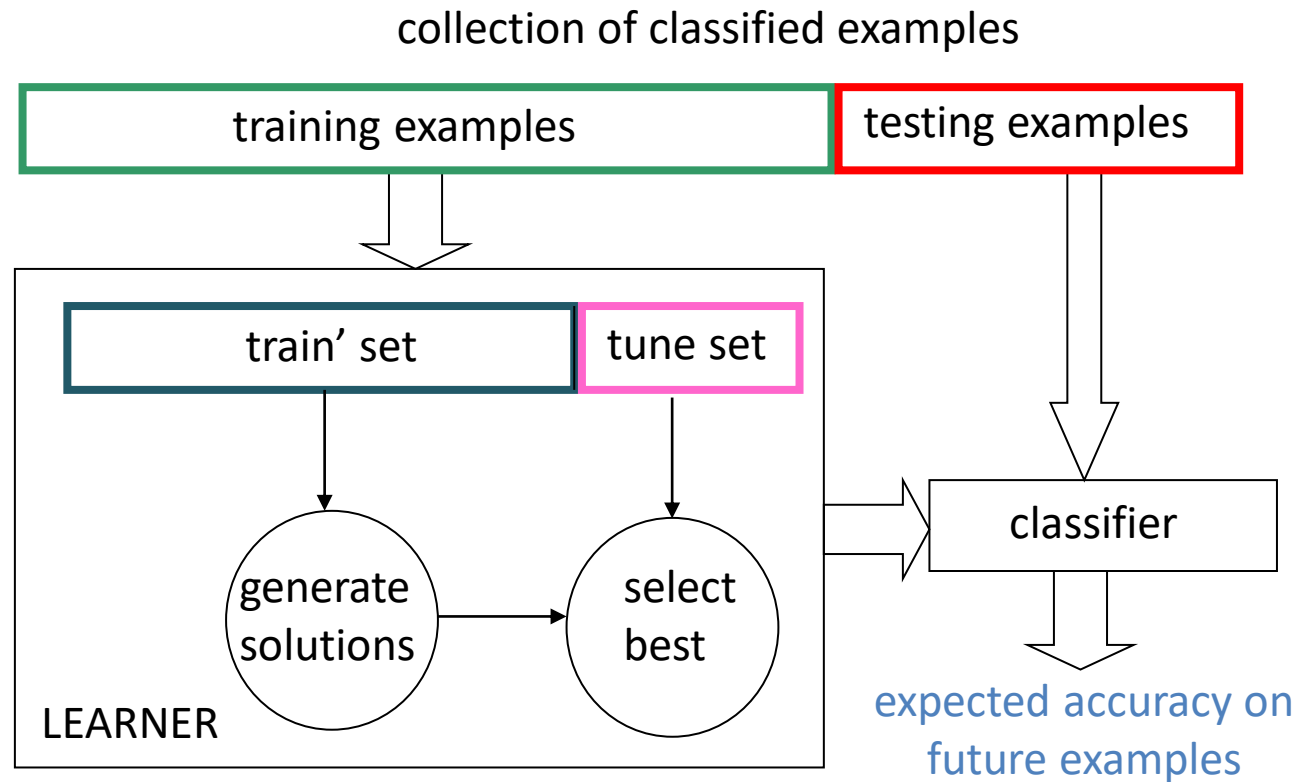
- 1) Start with a dataset of labeled examples
- 2) Randomly partition into  $N$  groups
- 3a)  $N$  times, combine  $N - 1$  groups into a train set
- 3b) Provide **training set** to learning system
- 3c) Measure accuracy on “left out” group (the **test set**)



Called  **$N$ -fold cross validation**

- Often, an ML system has to choose when to stop learning, select among alternative answers, etc.
- One wants the model that produces the highest accuracy on **future** examples (“overfitting avoidance”)
- It is a “**cheat**” to look at the **test** set while still learning
- Better method
  - Set aside part of the training set
  - Measure performance on this validation data to estimate future performance for a given set of hyperparameters
  - Use best hyperparameter settings, train with **all** training data (except **test** set) to estimate future performance on **new** examples

# A typical Learning system



Statistical techniques such as 10-fold cross validation and *t*-tests are used to get meaningful results

# Multiple Tuning sets



- Using a **single** tuning set can be unreliable predictor, plus some data “wasted”
  - 1) For each possible set of hyperparameters
    - a) Divide training data into **train** and **valid.** sets, using ***N*-fold cross validation**
    - b) Score this set of hyperparameter values: average **valid.** set accuracy over the *N* folds
  - 2) Use **best** set of hyperparameter settings and **all** (train + valid.) examples
  - 3) Apply resulting model to **test** set





# EVALUATING ML MODELS

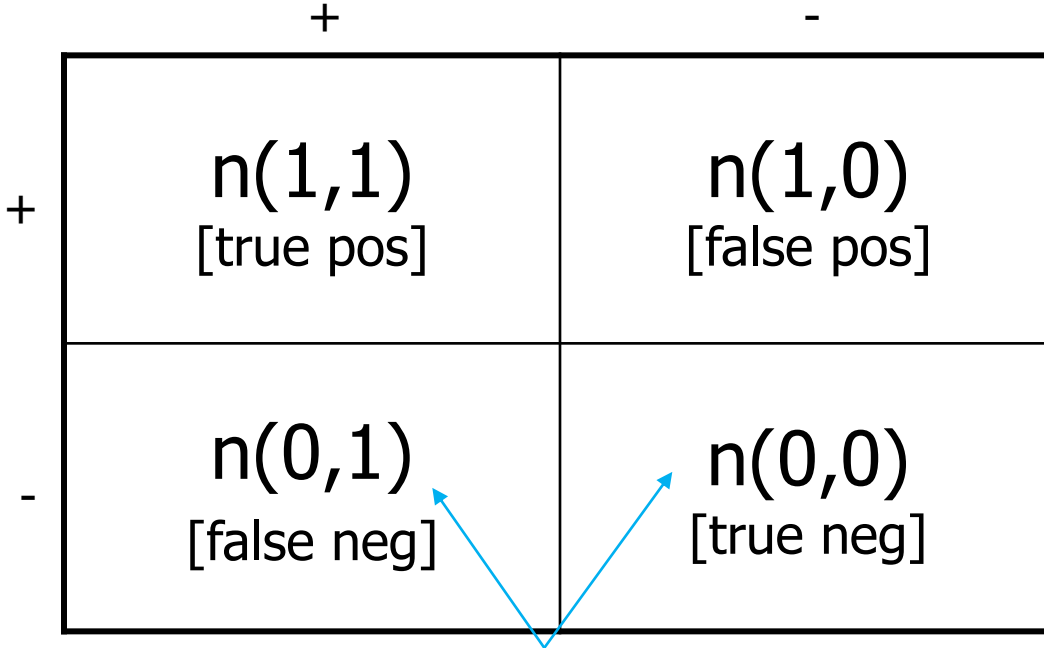
# Contingency Tables



(special case of 'confusion matrices')

		<b>True Answer</b>	
		+	-
<b>Algorithm Answer</b>	+	$n(1,1)$ [true pos]	$n(1,0)$ [false pos]
	-	$n(0,1)$ [false neg]	$n(0,0)$ [true neg]

Counts of occurrences




**True Positive Rate**  
(TPR)  $= n(1,1) / ( n(1,1) + n(0,1) )$   
 $=$  correctly categorized +’s / total positives  
 $\sim P(\text{algo outputs } + \mid + \text{ is correct})$

**False Positive Rate**  
(FPR)  $= n(1,0) / ( n(1,0) + n(0,0) )$   
 $=$  incorrectly categorized –’s / total neg’s  
 $\sim P(\text{algo outputs } + \mid - \text{ is correct})$

Can similarly define False Negative Rate and True Negative Rate

- ROC: *Receiver Operating Characteristics*
- Started for radar research during WWII
- Judging algorithms on accuracy alone may not be good enough when **getting a positive wrong** costs more than **getting a negative wrong** (or vice versa)
  - e.g., medical tests for serious diseases
  - e.g., a movie-recommender system

# ROC Curves Graphically



Different algorithms can work better in different parts of ROC space. This depends on cost of false + vs false -

# Creating an ROC Curve

---

## The Standard Approach:

- You need an ML algorithm that outputs **NUMERIC** results such as `prob(example is +)`
- You can use ensemble methods to get this from a model that only provides Boolean outputs
  - e.g., have 100 models vote & count votes

# Alg. for Creating ROC Curves

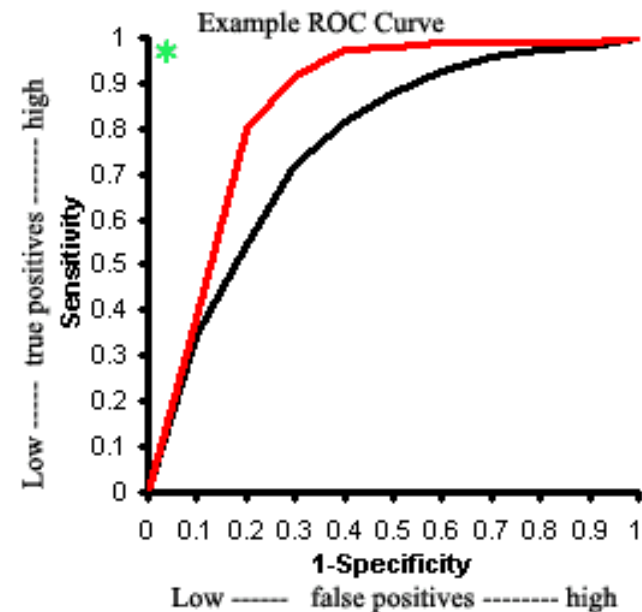


Step 1: Sort predictions on test set

Step 2: Locate a *threshold* between examples with opposite categories

Step 3: Compute TPR & FPR for each threshold of Step 2

Step 4: Connect the dots

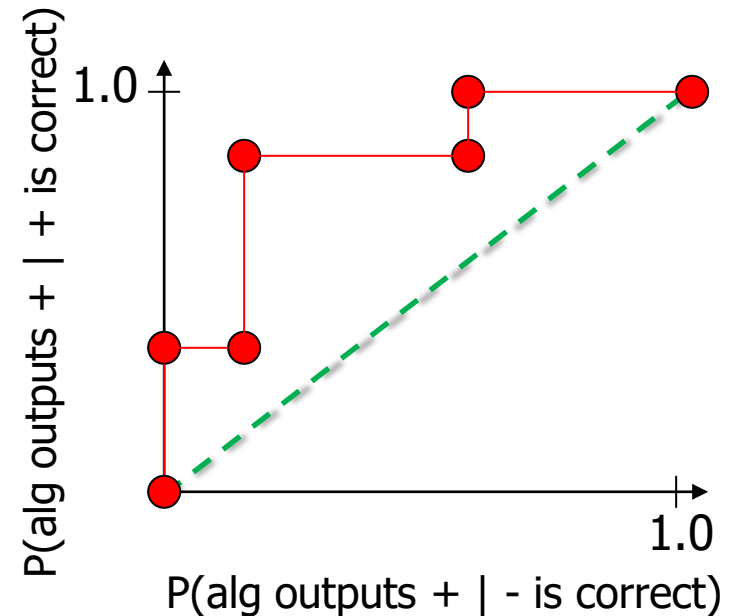


# Plotting ROC Curves - Example



ML Algo Output (Sorted)      Correct Category

Ex 9	.99		+
Ex 7	.98	TPR=(2/5), FPR=(0/5)	+
Ex 1	.72	TPR=(2/5), FPR=(1/5)	-
Ex 2	.70		+
Ex 6	.65	TPR=(4/5), FPR=(1/5)	+
Ex 10	.51		-
Ex 3	.39	TPR=(4/5), FPR=(3/5)	-
Ex 5	.24	TPR=(5/5), FPR=(3/5)	+
Ex 4	.11		-
Ex 8	.01	TPR=(5/5), FPR=(5/5)	-




Algorithm predicts + if its output is  $\geq 0$



# Area Under ROC Curve



- A common metric for experiments is to numerically integrate the ROC Curve
  - Usually called AUC
  - Probability that ML alg. will “rank” a randomly chosen positive instance higher than a randomly chosen negative one
  - Can summarize the curve **too much** in practice



# Asymmetric Error Costs



- Assume that  $\text{cost}(FP) \neq \text{cost}(FN)$
- You would like to pick a threshold that minimizes

$$\begin{aligned} E(\text{total cost}) \\ = \text{cost}(FP) \times \text{pr}(FP) \times (\# \text{ of neg ex's}) + \\ \text{cost}(FN) \times \text{pr}(FN) \times (\# \text{ of pos ex's}) \end{aligned}$$

- You could also have (maybe negative) costs for TP and TN (assumed zero in above)

- One strength of ROC curves is that they are a good way to deal with **skewed** data ( $|+| \gg |-|$ ) since the axes are fractions (rates) independent of the # of examples
- You must be careful though!
  - Low FPR \* (many negative ex) = **sizable number of FP**
  - Possibly more than # of TP

# Precision vs. Recall




- Think about search engines...
- **Precision** = 
$$\frac{\text{\# of relevant items retrieved}}{\text{\# of total items retrieved}}$$
$$= n(1,1) / ( n(1,1) + n(1,0) )$$
$$\cong P(\text{is pos} \mid \text{called pos})$$
- **Recall** = 
$$\frac{\text{\# of relevant items retrieved}}{\text{\# of relevant items that exist}}$$
$$= n(1,1) / ( n(1,1) + n(0,1) ) = \underline{\text{TPR}}$$
$$\cong P(\text{called pos} \mid \text{is pos})$$
- Notice that  $n(0,0)$  is not used in either formula  
Therefore you get no credit for filtering out irrelevant items

# ROC vs. Precision-Recall



You can get very different visual results  
on the same data!




The reason for this is that there may be lots of – ex's  
(e.g., might need to include 100 neg's to get 1 more pos)

- In most learning algorithms, we can specify a threshold for making a rejection decision
- Probabilistic classifiers: adjust cost of rejecting versus cost of FP and FN
- Decision-boundary method: if a test point  $x$  is within  $\theta$  of the decision boundary, then reject
  - Equivalent to requiring that the “activation” of the best class is larger than the second-best class by at least  $\theta$

# Rejection Curves



- Vary  $\theta$  and plot fraction correct versus fraction rejected



# The F1 Measure



- Figure of merit that combines precision and recall

$$F_1 = 2 \cdot \frac{P \cdot R}{P + R}$$

where  $P$  = precision;  $R$  = recall. This is twice the harmonic mean of  $P$  and  $R$ .

- We can plot  $F_1$  as a function of the classification threshold  $\theta$