



# Nearest Neighbor Methods

Rishabh Iyer

University of Texas at Dallas

Based on the slides of Vibhav Gogate, Nick Rouzzi, David Sontag and few other sources

# Nearest Neighbor Methods




- Learning
  - Store all training examples

$$\mathcal{T} = \underbrace{\{(x_1, y_1), \dots, (x_M, y_M)\}}_{M \text{ Train Data points}}$$


- Classifying a new point  $x'$ 
  - Find the training example  $(x^{(i)}, y^{(i)})$  such that  $x^{(i)}$  is closest (for some notion of close) to  $x'$
  - Classify  $x'$  with the label  $y^{(i)}$

$x' \rightarrow$  Test point.


# Nearest Neighbor Methods




# Nearest Neighbor Methods



# Nearest Neighbor Methods



# Nearest Neighbor Methods




# Nearest Neighbor Methods



1-NN : Closest.


K-NN :  $k$  Closest points

$x^1 : \{y^1, \dots, y^K\}$



$k$ -nearest neighbor methods look at the  $k$  closest points in the training set and take a majority vote  
(should choose  $k$  to be odd)


# Nearest Neighbor Methods



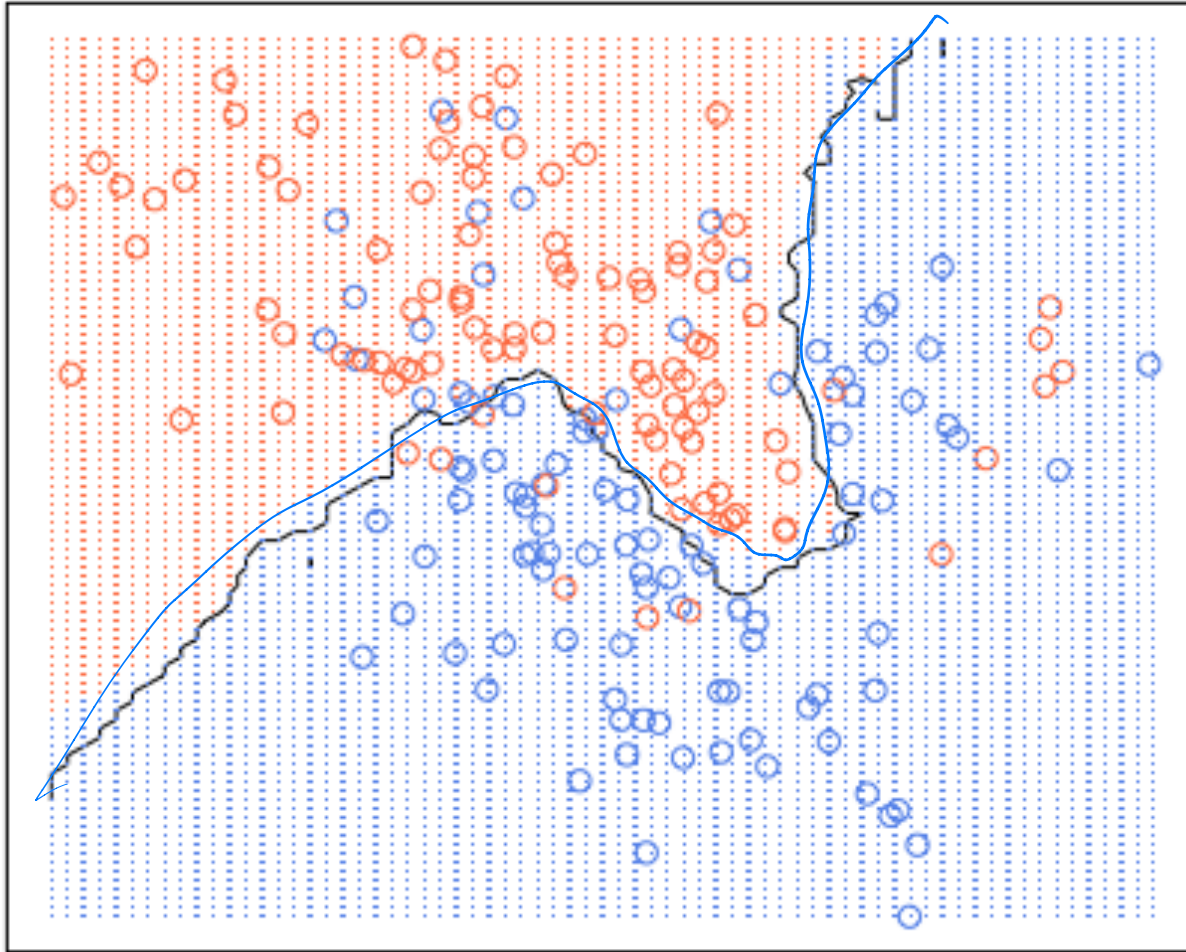
$k$ -nearest neighbor methods look at the  $k$  closest points in the training set and take a majority vote (should choose  $k$  to be odd)



# 1-NN Example


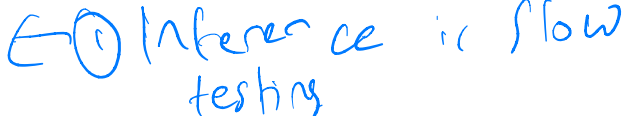



# 20-NN Example



# Nearest Neighbor Methods



- Applies to data sets with points in  $\mathbb{R}^d$ 
  - Best for large data sets with only a few ( $< 20$ ) attributes
- Advantages
  - Learning is easy  *no training*
  - Can learn complicated decision boundaries
- Disadvantages
  - Classification is slow (need to keep the entire training set around)  $M \sim |M|$   *Inference is slow testing*
  - Easily fooled by irrelevant attributes  *High memory*

- How to choose the right measure of closeness?
  - Euclidean distance is popular, but many other possibilities

- How to pick  $k$ ?
  - Too small and the estimates are noisy, too large and the accuracy suffers

$$x', x^{(i)} \in \mathbb{R}^d$$

- What if the nearest neighbor is really far away?

$$\begin{array}{ccc} \parallel x' - x^{(i)} \parallel_2^2 & = & \sum_{j=1}^d (x'_j - x_j^{(i)})^2 \\ \uparrow & & \uparrow \\ \text{test} & & \text{train} \end{array}$$

# Choosing the Distance



- Euclidean distance makes sense when each of the features is roughly on the same scale
  - If the features are very different (e.g., <sup>5</sup>height and <sup>20-30</sup>age), then Euclidean distance makes less sense as height would be less significant than age simply because age has a larger range of possible values
  - To correct for this, feature vectors are often recentered around their means and scaled by the standard deviation over the training set

# Normalization

[Mean - Variance Normalization]



- Sample mean

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x^{(i)}$$

- Sample variance (biased)

$$\hat{\sigma}_k^2 = \frac{1}{n} \sum_{i=1}^n (x_k^{(i)} - \bar{x}_k)^2$$

$x_k$   
 $\downarrow$   
 $x_k = \frac{x_k - \bar{x}_k}{\sigma_k}$

# Min-Max Normalization

$$x_k^{\max} = \max_{i=1:n} x_k$$


$$x_k^{\min} = \min_{i=1:n} x_k \quad (i)$$

$$x_k'' = \frac{x_k - x_k^{\min}}{x_k^{\max} - x_k^{\min}} \in [0, 1]$$

# Irrelevant Attributes



Consider the nearest neighbor problem in one dimension






# Irrelevant Attributes





Now, add a new attribute that is just random noise...



$$1 \rightarrow 2 \rightarrow 3$$

$$d_{12} \leq d_{13}$$



# K-Dimensional Trees



*inference / testing*

- In order to do classification, we can compute the distances between all points in the training set and the point we are trying to classify
- With <sup>1M</sup> $m$  data points in <sup>100</sup> $n$ -dimensional space, this takes  $O(mn)$  time for Euclidean distance
- It is possible to do better if we do some preprocessing on the training data


*Test:  $n_t$  ~ 100M (slow?)*  
*Prediction ~ MS*

- k-d trees provide a data structure that can help simplify the classification task by constructing a tree that partitions the search space
  - Starting with the entire training set, choose some dimension,  $i$
  - Select an element of the training data whose  $i^{th}$  dimension has the median value among all elements of the training set
  - Divide the training set into two pieces: depending on whether their  $i^{th}$  attribute is smaller or larger than the median
  - Repeat this partitioning process on each of the two new pieces separately

# K-Dimensional Trees




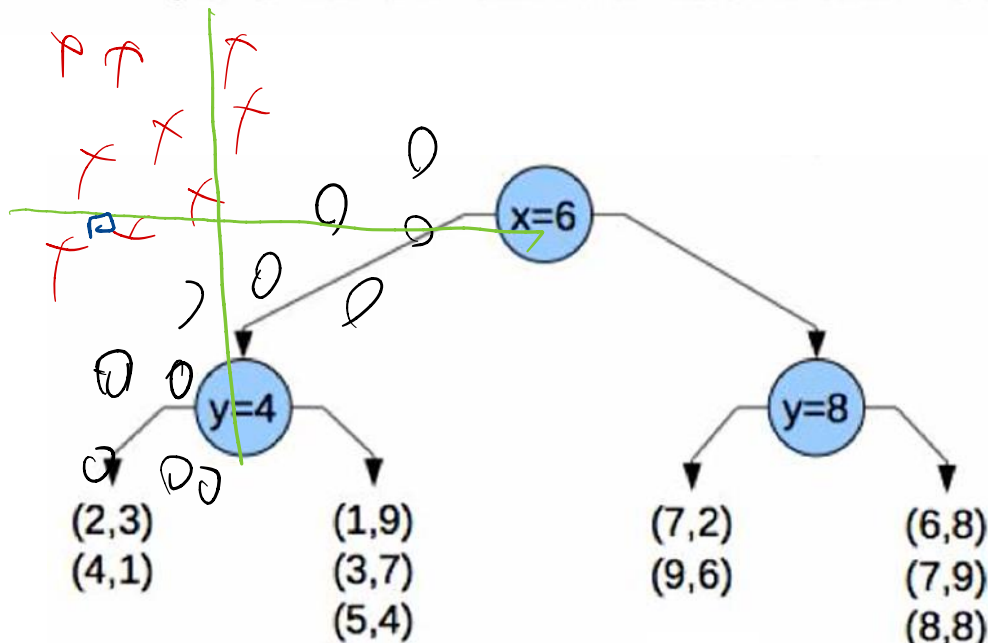
- Building a K-D tree from training data:
  - $\{(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)\}$
  - pick random dimension, find median, split data, repeat



# K-Dimensional Trees



- Building a K-D tree from training data:
  - $\{(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)\}$
  - pick random dimension, find median, split data, repeat





# K-Dimensional Trees: Inference



- Find NNs for new point (7,4)  $\leftarrow$  Test
  - find region containing (7,4)
  - compare to all points in region







Approximate NN



- By design, the constructed k-d tree is “bushy”
  - The idea is that if new points to classify are evenly distributed throughout the space, then the expected (amortized) cost of classification is approximately  $O(d \log n)$  operations
- Summary
  - k-NN is fast and easy to implement
  - No training required
  - Can be good in practice (where applicable)



K-mediod.



$$10 = m.$$

Color: { red, green, blue, ... }

→ [ 1 0 0 ... 0 ]

1-hot encoding.

[ 0 0 1 0 0 ... 0 ]

